

▼ CSE527 Homework 4

Due date: 23:59 on Nov. 5, 2019 (Tuesday)

In this semester, we will use Google Colab for the assignments, which allows us to utilize resources of local machines such as GPUs. You will need to use your Stony Brook (*.stonybrook.edu) account for course results.

Google Colab Tutorial

Go to <https://colab.research.google.com/notebooks/>, you will see a tutorial named "Welcome to Colab basics of using google colab."

Settings used for assignments: **Edit -> Notebook Settings -> Runtime Type (Python 3)**.

Description

This project is an introduction to deep learning tools for computer vision. You will design and train deep recognition using [PyTorch](#). You can visualize the structure of the network with [mNeuron] (<http://vision>

Remember Homework 3: Scene recognition with bag of words. You worked hard to design a bag of features to 70% accuracy (most likely) on 16-way scene classification. We're going to attack the same task with Training from scratch won't work quite as well as homework 3 due to the insufficient amount of data, but much better than homework 3.

In Problem 1 of the project you will train a deep convolutional network from scratch to recognize scenes to load data and display them. You will need to define a simple network architecture and add jittering, increase recognition accuracy to 50, 60, or perhaps 70%. Unfortunately, we only have 2,400 training examples to train a network from scratch which outperforms hand-crafted features

For Problem 2 you will instead fine-tune a pre-trained deep network to achieve about 85% accuracy on AlexNet network which was not trained to recognize scenes at all.

These two approaches represent the most common approaches to recognition problems in computer vision: training from scratch if you have enough data (it's not always obvious whether or not you do), and if you cannot then fine-tuning.

There are 2 problems in this homework with a total of 110 points including 10 bonus points. Be sure to complete all problems as they are important. For the problems requiring text descriptions, you might want to add a markdown block

Dataset

Save the [dataset\(click me\)](#) into your working folder in your Google Drive for this homework.

Under your root folder, there should be a folder named "data" (i.e. XXX/Surname_Givenname_SBUID/data)

the data subfolder before submitting on blackboard due to size limit. There should be only one .ipynb Surname_Givenname_SBUID.

Some Tutorials (PyTorch)

- You will be using PyTorch for deep learning toolbox (follow the [link](#) for installation).
- For PyTorch beginners, please read this [tutorial](#) before doing your homework.
- Feel free to study more tutorials at <http://pytorch.org/tutorials/>.
- Find cool visualization here at <http://playground.tensorflow.org>.

Starter Code

In the starter code, you are provided with a function that loads data into minibatches for training and t

```
# import packages here
import cv2
import numpy as np
import matplotlib.pyplot as plt
import glob
import random
import time

import torch
import torchvision
import torchvision.transforms as transforms

from torch.autograd import Variable
import torch.nn as nn
import torch.nn.functional as F

import pickle
import torch.optim as optim

# Mount your google drive where you've saved your assignment folder
from google.colab import drive
drive.mount('/content/gdrive')

📁 Drive already mounted at /content/gdrive; to attempt to forcibly remount, call dr

# Set your working directory (in your google drive)
# Note that 'gdrive/My Drive/Y2019Fall/CSE-527-Intro-To-Computer-Vision/hw4' is just a
# change it to your specific homework directory.
cd '/content/gdrive/My Drive/ComputerVision_Fall2019/Md_Arif_112669645_hw4/'
```

```
↳ /content/drive/My Drive/ComputerVision_Fall2019/Md_Arif_112669645_hw4
```

```
print(torch.__version__)
print(torch.cuda.is_available())
print(torch.version.cuda)
```

```
↳ 1.3.0+cu100
True
10.0.130
```

```
# =====
# Load Training Data and Testing Data
# =====
```

```
class_names = [name[13:] for name in glob.glob('./data/train/*')]
class_names = dict(zip(range(len(class_names)), class_names))
print("class_names: %s " % class_names)
n_train_samples = 150
n_test_samples = 50
```

```
def img_norm(img):
    # Write your code here
    # normalize img pixels to [-1, 1]
    norm_img= np.float32(cv2.normalize(img,None,0,1,cv2.NORM_MINMAX))
    norm_img-=0.5
    return norm_img*2
```

```
def load_dataset(path, img_size, num_per_class=-1, batch_num=1, shuffle=False, augment
                rotate_90=False, zero_centered=False):
```

```
    data = []
    labels = []
```

```
    if is_color:
        channel_num = 3
    else:
        channel_num = 1
```

```
    # read images and resizing
    for id, class_name in class_names.items():
        print("Loading images from class: %s" % id)
        img_path_class = glob.glob(path + class_name + '/*.jpg')
        if num_per_class > 0:
            img_path_class = img_path_class[:num_per_class]
        labels.extend([id]*len(img_path_class))
        for filename in img_path_class:
            if is_color:
                img = cv2.imread(filename)
            else:
                img = cv2.imread(filename, 0)
```

```
        # resize the image
```

```

img = cv2.resize(img, img_size, cv2.INTER_LINEAR)

if is_color:
    img = np.transpose(img, [2, 0, 1])

# norm pixel values to [-1, 1]
data.append(img_norm(img))

# Write your Data Augmentation code here
# mirroring

if augment:
    for id, class_name in class_names.items():
        print("Loading images from class: %s" % id)
        img_path_class = glob.glob(path + class_name + '/*.jpg')
        if num_per_class > 0:
            img_path_class = img_path_class[:num_per_class]
        labels.extend([id]*len(img_path_class))
        for filename in img_path_class:
            if is_color:
                img = cv2.imread(filename)
                img = img[:, ::-1] ## Taking Mirror Image
            else:
                img = cv2.imread(filename, 0)
                img = img[:, ::-1]

            # resize the image
            img = cv2.resize(img, img_size, cv2.INTER_LINEAR)

            if is_color:
                img = np.transpose(img, [2, 0, 1])

            # norm pixel values to [-1, 1]
            data.append(img_norm(img))
    if rotate_90:
        img_rotate_clock = [np.rot90(img, axes=(-2,-1)) for img in data]
        data.extend(img_rotate_clock)
        labels.extend(labels)

# randomly permute (this step is important for training)
if shuffle:
    bundle = list(zip(data, labels))
    random.shuffle(bundle)
    data, labels = zip(*bundle)

# divide data into minibatches of TorchTensors
if batch_num > 1:
    batch_data = []
    batch_labels = []

    print(len(data))
    print(batch_num)

```

```

for i in range(int(len(data) / batch_num)):
    minibatch_d = data[i*batch_num: (i+1)*batch_num]
    minibatch_d = np.reshape(minibatch_d, (batch_num, channel_num, img_size[0],
    batch_data.append(torch.from_numpy(minibatch_d))

    minibatch_l = labels[i*batch_num: (i+1)*batch_num]
    batch_labels.append(torch.LongTensor(minibatch_l))
data, labels = batch_data, batch_labels

return zip(batch_data, batch_labels)

```

```

class_names: {0: 'TallBuilding', 1: 'InsideCity', 2: 'Mountain', 3: 'LivingRoom',

```

```

# load data into size (64, 64)

```

```

start_time = time.time()

```

```

img_size = (64, 64)

```

```

batch_num = 50 # training sample number per batch

```

```

# load training dataset

```

```

trainloader_small = list(load_dataset('./data/train/', img_size, batch_num=batch_num,
                                     augment=True, zero_centered=True))

```

```

train_num = len(trainloader_small)

```

```

print("Finish loading %d minibatches(=%d) of training samples." % (train_num, batch_num))

```

```

# load testing dataset

```

```

testloader_small = list(load_dataset('./data/test/', img_size, num_per_class=50, batch_num=batch_num))
test_num = len(testloader_small)

```

```

print("Finish loading %d minibatches(=%d) of testing samples." % (test_num, batch_num))

```

```

print('\033[1m' + 'Time Taken to load the Training and Test Data Set: ' + '--- %s minutes' % (time.time() - start_time))

```


```


```

```
Loading images from class: 0
Loading images from class: 1
Loading images from class: 2
Loading images from class: 3
Loading images from class: 4
Loading images from class: 5
Loading images from class: 6
Loading images from class: 7
Loading images from class: 8
Loading images from class: 9
Loading images from class: 10
Loading images from class: 11
Loading images from class: 12
Loading images from class: 13
Loading images from class: 14
Loading images from class: 15
Loading images from class: 0
Loading images from class: 1
Loading images from class: 2
Loading images from class: 3
Loading images from class: 4
Loading images from class: 5
Loading images from class: 6
Loading images from class: 7
Loading images from class: 8
Loading images from class: 9
Loading images from class: 10
Loading images from class: 11
Loading images from class: 12
Loading images from class: 13
Loading images from class: 14
Loading images from class: 15
4800
50
Finish loading 96 minibatches(=50) of training samples.
Loading images from class: 0
Loading images from class: 1
Loading images from class: 2
Loading images from class: 3
Loading images from class: 4
Loading images from class: 5
Loading images from class: 6
Loading images from class: 7
Loading images from class: 8
Loading images from class: 9
Loading images from class: 10
Loading images from class: 11
Loading images from class: 12
Loading images from class: 13
Loading images from class: 14
Loading images from class: 15
400
50
Finish loading 8 minibatches(=50) of testing samples.
Time Taken to load the Training and Test Data Set: --- 2.86 minutes ---
```

```
# show some images
def imshow(img):
    img = img / 2 + 0.5      # unnormalize
    npimg = img.numpy()
    if len(npimg.shape) > 2:
        npimg = np.transpose(img, [1, 2, 0])
    plt.figure
    plt.imshow(npimg, 'gray')
    plt.show()
img, label = trainloader_small[0][0][11][0], trainloader_small[0][1][11]
label = int(np.array(label))
print(class_names[label])
imshow(img)
```

📄 Flower



```
# # As loading the data from the source for the first time is time consuming, so you c
# # Save intermediate image data into disk
file = open('trainloader_small.pkl','wb')
pickle.dump(trainloader_small, file)

file = open('testloader_small.pkl','wb')
pickle.dump(testloader_small, file)
file.close()

# Load intermediate image data from disk
file = open('trainloader_small.pkl', 'rb')
trainloader_small = pickle.load(file)
file.close()

file = open('testloader_small.pkl', 'rb')
testloader_small = pickle.load(file)
file.close()

print('\033[1m' + 'Training Samples: ' + str(len(trainloader_small)*50)) # Verify num
print('\033[1m' + 'Test Samples: ' + str(len(testloader_small)*50)) # Verify number
```

```

↳ Training Samples: 4800
Test Samples: 400

```

```

print('\033[1m' + 'Total Training Batches: ' + str(len(trainloader_small)))
print('\033[1m' + 'Total Test Batches: ' + str(len(testloader_small)))

```

```

↳ Total Training Batches: 96
Total Test Batches: 8

```

```

## Basic Data Sanity Checks
sample = next(iter(trainloader_small))
print(len(sample))
print(type(sample))
smpl_img, smpl_lbl= sample

i=10
print(smpl_img[i].shape)
print(smpl_lbl[i].shape)

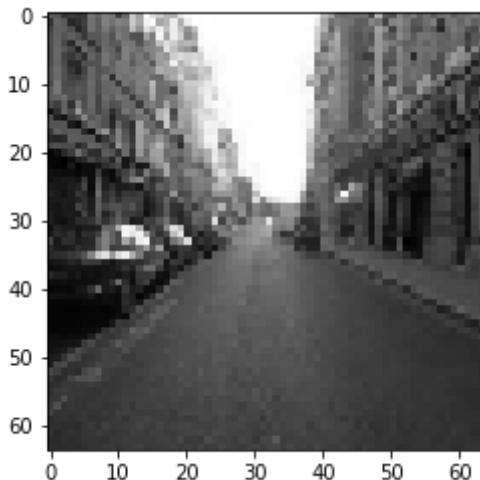
plt.imshow(smpl_img[0].squeeze(),cmap= 'gray')
print('label',smpl_lbl[i] )

```

```

↳ 2
<class 'tuple'>
torch.Size([1, 64, 64])
torch.Size([])
label tensor(0)

```



▼ Problem 1: Training a Network From Scratch

{Part 1: 35 points} Gone are the days of hand designed features. Now we have end-to-end learning in \ learned for our data to maximize our objective (in this case, 16-way classification accuracy). Instead c scenes with... 25% accuracy. OK, that didn't work at all. Try to boost the accuracy by doing the followir

Data Augmentation: We don't have enough training data, let's augment the training data. If you left-right mirror images, the categories never change. A kitchen doesn't become a forest when mirrored. This isn't true in all domains, so you can't "jitter" digit recognition training data in the same way. But we can synthetically increase our training data by mirroring training images during the learning process.

After you implement mirroring, you should notice that your training error doesn't drop as quickly. That's because the network isn't overfitting to the 2,400 original training images as much (because it sees 4,800 training images, good as 4,800 truly independent samples). Because the training and test errors fall more slowly, you may want to try modifying the learning rate. You should see a roughly 10% increase in accuracy by adding mirroring as data augmentation for this part.

You can try more elaborate forms of jittering -- zooming in a random amount, rotating a random amount. If not required, you might want to try these in the bonus part.

Data Normalization: The images aren't zero-centered. One simple trick which can help a lot is to subtract the mean from the training images (since the test/validation set is also from the same distribution, this will arguably be more proper to only compute the mean from the training images (since the test/validation set won't make much of a difference. After doing this you should see another 15% or so increase in accuracy).

Network Regularization: Add dropout layer. If you train your network (especially for more than the default number of epochs), the training error can decrease to zero while the val top1 error hovers at 40% to 50%. The network has learned weights that generalize to the training data, but those weights don't generalize to held out test data. The best regularization would be to use dropout. Instead we will use dropout regularization.

What does dropout regularization do? It randomly turns off network connections at training time to force each layer from relying too strongly on a single unit in the previous layer. Dropout regularization can be interpreted as training "thinned" versions of your network. At test, all connections are restored which is analogous to taking an average over the "thinned" networks. You can see a more complete discussion of dropout regularization in this [paper](#).

The dropout layer has only one free parameter -- the dropout rate -- the proportion of connections that should be dropped. Insert a dropout layer between your convolutional layers. In particular, insert it directly after the first convolutional layer. Your test accuracy should increase by another 10%. Your train accuracy should decrease much more slowly. This makes life much harder for the training algorithm by cutting out connections randomly.

If you increase the number of training epochs (and maybe decrease the learning rate) you should be able to reach higher accuracy. In this part, you are **required** to add dropout layer to your network.

Please give detailed descriptions of your network layout in the following format:

Data augmentation: [descriptions]

Data normalization: [descriptions]

Layer 1: [layer_type]: [Parameters]

Layer 2: [layer_type]: [Parameters]

...

Then report the final accuracy on test set and time consumed for training and testing separately.

{Part 2: 15 points} Try **three techniques** taught in the class to increase the accuracy of your model. Such as randomly rotating training images, adding batch normalization, different activation functions (e.g., sigmoid modification). Note that too many layers can do you no good due to insufficient training data. Clearly document the increase/decrease for each of the three techniques.

```
# =====
#         Define Network Architecture
# =====

class CNN_Ver1(nn.Module):
    def __init__(self, input_channels=1, num_classes=16):
        super(CNN_Ver1, self).__init__()
        self.layer1 = nn.Sequential()
        self.layer1.add_module("Conv1", nn.Conv2d(in_channels=input_channels, out_channels=16, kernel_size=3, padding=1))
        self.layer1.add_module("Relu1", nn.ReLU())
        self.layer1.add_module("MaxPool1", nn.MaxPool2d(kernel_size=2))

        self.layer2 = nn.Sequential()
        self.layer2.add_module("Conv2", nn.Conv2d(in_channels=16, out_channels=32, kernel_size=3, padding=1))
        self.layer2.add_module("Relu2", nn.ReLU())
        self.layer2.add_module("MaxPool2", nn.MaxPool2d(kernel_size=2))

        self.fully_connected = nn.Linear(32 * 16 * 16, num_classes)

    def forward(self, x):
        x = self.layer1(x)
        x = self.layer2(x)
        x = x.view(x.size(0), -1)
        x = self.fully_connected(x)
        return x

# defining the model
model_ver1 = CNN_Ver1()

# defining the optimizer
optimizer = optim.Adam(model_ver1.parameters(), lr=0.001)

# defining the loss function
criterion = nn.CrossEntropyLoss()
# checking if GPU is available
if torch.cuda.is_available():
    model_ver1 = model_ver1.cuda()
    criterion = criterion.cuda()

print(model_ver1)
```



```

CNN_Ver1(
    (layer1): Sequential(
      (Conv1): Conv2d(1, 16, kernel_size=(5, 5), stride=(1, 1), padding=(2, 2))
      (Relu1): ReLU()
      (MaxPool1): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mo
    )
    (layer2): Sequential(
      (Conv2): Conv2d(16, 32, kernel_size=(5, 5), stride=(1, 1), padding=(2, 2))
      (Relu2): ReLU()
      (MaxPool2): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mo
    )
    (fully_connected): Linear(in_features=8192, out_features=16, bias=True)
  )

```

```

# =====
#           Optimize/Train Network
# =====
## Defining Training Framework

```

```

def train_model(train_set, epoch,model_version,print_mini_batch_loss= False, print_epo
    model_version.train()
    if print_epoch_loss:
        print('\033[1m' + 'Performance Statistics for Epoch Number :' + str(epoch + 1)
    train_running_loss = 0.0
    train_total_loss= 0.0
    for batch_num, data in enumerate(train_set, 0):
        # getting the inputs; data is a tensor tuple list of [inputs, labels]
        x_train, y_train = data
        x_train = x_train.float()

        # converting the data into GPU format
        if torch.cuda.is_available():
            x_train = x_train.cuda()
            y_train = y_train.cuda()

        # zero the parameter gradients
        # prediction for training set
        # computing the training loss
        # computing the updated weights of all the model parameters

        optimizer.zero_grad()
        output_train = model_version(x_train)
        loss_train = criterion(output_train, y_train)
        loss_train.backward()
        optimizer.step()

        # print statistics
        train_running_loss += loss_train.item()
        train_total_loss += loss_train.item()

        if batch_num % 20 == 19:      # print every 20 mini-batches

```

```

        if print_mini_batch_loss:
            print('\033[1m' + '[Batch NO.: %d, Mini Batch: %5d] loss: %.3f' %
                  (epoch + 1, batch_num + 1, train_running_loss / 20))
            train_running_loss = 0.0
        train_losses.append(train_total_loss/(batch_num + 1))
    if print_epoch_loss:
        print('\033[1m' + '[Avg Loss for the Epoch NO.: %d, Loss: %.3f]' % (epoch + 1,
        print('\n')

```

```

start_time = time.time()
# defining the number of epochs
n_epochs = 25
# empty list to store training losses
train_losses = []
# loop over the dataset multiple times
for epoch in range(n_epochs):
    train_model(trainloader_small,epoch,model_ver1, False, True)

print('Finished Training')
print('\033[1m' + 'Time Taken to Train the Model: ' + '--- %s mins ---' % round((time.

```



Performance Statistics for Epoch Number :1
[Avg Loss for the Epoch NO.: 1, Loss: 2.361]

Performance Statistics for Epoch Number :2
[Avg Loss for the Epoch NO.: 2, Loss: 1.632]

Performance Statistics for Epoch Number :3
[Avg Loss for the Epoch NO.: 3, Loss: 1.301]

Performance Statistics for Epoch Number :4
[Avg Loss for the Epoch NO.: 4, Loss: 1.092]

Performance Statistics for Epoch Number :5
[Avg Loss for the Epoch NO.: 5, Loss: 0.935]

Performance Statistics for Epoch Number :6
[Avg Loss for the Epoch NO.: 6, Loss: 0.802]

Performance Statistics for Epoch Number :7
[Avg Loss for the Epoch NO.: 7, Loss: 0.665]

Performance Statistics for Epoch Number :8
[Avg Loss for the Epoch NO.: 8, Loss: 0.544]

Performance Statistics for Epoch Number :9
[Avg Loss for the Epoch NO.: 9, Loss: 0.443]

Performance Statistics for Epoch Number :10
[Avg Loss for the Epoch NO.: 10, Loss: 0.410]

Performance Statistics for Epoch Number :11
[Avg Loss for the Epoch NO.: 11, Loss: 0.388]

Performance Statistics for Epoch Number :12
[Avg Loss for the Epoch NO.: 12, Loss: 0.310]

Performance Statistics for Epoch Number :13
[Avg Loss for the Epoch NO.: 13, Loss: 0.228]

Performance Statistics for Epoch Number :14
[Avg Loss for the Epoch NO.: 14, Loss: 0.172]

Performance Statistics for Epoch Number :15

[Avg Loss for the Epoch NO.: 15, Loss: 0.125]

Performance Statistics for Epoch Number :16

[Avg Loss for the Epoch NO.: 16, Loss: 0.122]

Performance Statistics for Epoch Number :17

[Avg Loss for the Epoch NO.: 17, Loss: 0.092]

Performance Statistics for Epoch Number :18

[Avg Loss for the Epoch NO.: 18, Loss: 0.066]

Performance Statistics for Epoch Number :19

[Avg Loss for the Epoch NO.: 19, Loss: 0.051]

Performance Statistics for Epoch Number :20

[Avg Loss for the Epoch NO.: 20, Loss: 0.034]

Performance Statistics for Epoch Number :21

[Avg Loss for the Epoch NO.: 21, Loss: 0.019]

Performance Statistics for Epoch Number :22

[Avg Loss for the Epoch NO.: 22, Loss: 0.013]

Performance Statistics for Epoch Number :23

[Avg Loss for the Epoch NO.: 23, Loss: 0.009]

Performance Statistics for Epoch Number :24

[Avg Loss for the Epoch NO.: 24, Loss: 0.006]

Performance Statistics for Epoch Number :25

[Avg Loss for the Epoch NO.: 25, Loss: 0.004]

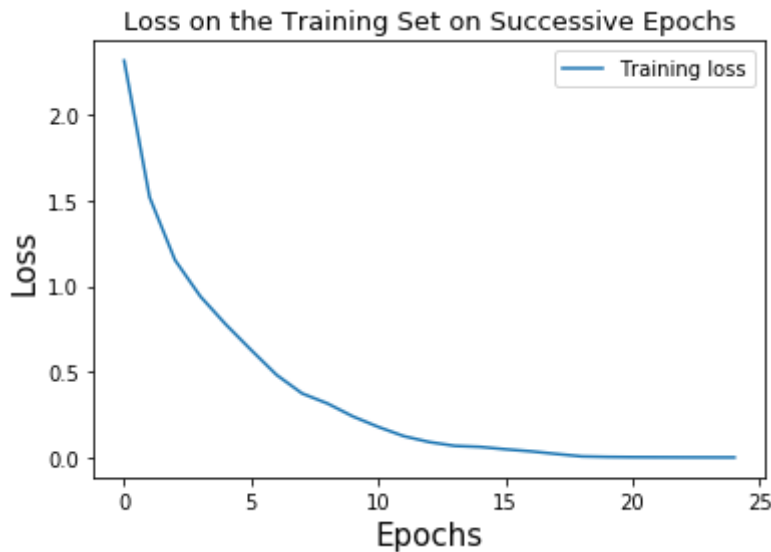
Finished Training

Time Taken to Train the Model: --- 0.47 mins ---

```
file = open('model_ver1.pkl','wb')
pickle.dump(model_ver1, file)
file.close()
```

```
file = open('model_ver1.pkl', 'rb')
model_ver1 = pickle.load(file)
file.close()
```

```
# plotting the training loss
plt.plot(train_losses, label='Training loss')
plt.title('Loss on the Training Set on Successive Epochs',size= 13)
plt.legend()
plt.xlabel('Epochs',size= 15)
plt.ylabel('Loss',size=15)
plt.show()
```



```
# =====
#           Evaluating Network
# =====
```

```
def test_model(test_data,model_version):
    start_time = time.time()
    # prediction and Accuracy for Test set
    correct = 0
    total = 0
    model_version.eval()
    for batch_num, data in enumerate(test_data, 0):
        x_test, y_test = data
        x_test = x_test.float()

        if torch.cuda.is_available():
            x_test = x_test.cuda()
            y_test = y_test.cuda()

        y_pred = model_version(x_test.cuda())
        _, predicted = torch.max(y_pred.data, 1)
        total += y_test.size(0)
        correct += (predicted == y_test).sum().item()

    print('\033[1m' + 'Accuracy of the network on the test images: %d %%' % (100 * cor
    print('\033[1m' + 'Time Taken for Evaluating Model on Test: ' + '--- %s seconds --
```

```
test_model(testloader_small,model_ver1)
```

```
↳ Accuracy of the network on the test images: 62 %
Time Taken for Evaluating Model on Test: --- 0.05 seconds ---
```

Part 2 : Let's Try to Increase Accuracy by Following Techniques

▼ Method1: Adding Batch Normalization:

```
class CNN_Ver2(nn.Module):
    def __init__(self, input_channels=1, num_classes=16):
        super(CNN_Ver2, self).__init__()
        self.layer1 = nn.Sequential()
        self.layer1.add_module("Conv1", nn.Conv2d(input_channels, out_channels=16, kernel_size=3, padding=1))
        self.layer1.add_module("BN1", nn.BatchNorm2d(num_features=16))
        self.layer1.add_module("Relu1", nn.ReLU())
        self.layer1.add_module("MaxPool1", nn.MaxPool2d(kernel_size=2))

        self.layer2 = nn.Sequential()
        self.layer2.add_module("Conv2", nn.Conv2d(16, out_channels=32, kernel_size=3, padding=1))
        self.layer2.add_module("BN2", nn.BatchNorm2d(num_features=32))
        self.layer2.add_module("Relu2", nn.ReLU())
        self.layer2.add_module("MaxPool2", nn.MaxPool2d(kernel_size=2))

        self.fully_connected = nn.Linear(32 * 16 * 16, num_classes)

    def forward(self, x):
        x = self.layer1(x)
        x = self.layer2(x)
        x = x.view(x.size(0), -1)
        x = self.fully_connected(x)
        return x

# defining the model
model_ver2 = CNN_Ver2()
optimizer = optim.Adam(model_ver2.parameters(), lr=0.001)
criterion = nn.CrossEntropyLoss()

if torch.cuda.is_available():
    model_ver2 = model_ver2.cuda()
    criterion = criterion.cuda()

print(model_ver2)
```

```
↳
```



```

CNN_Ver2(
  (layer1): Sequential(
    (Conv1): Conv2d(1, 16, kernel_size=(5, 5), stride=(1, 1), padding=(2, 2))
    (BN1): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True, track_running_st
    (Relu1): ReLU()
    (MaxPool1): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mo
  )
  (layer2): Sequential(
    (Conv2): Conv2d(16, 32, kernel_size=(5, 5), stride=(1, 1), padding=(2, 2))
    (BN2): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True, track_running_st
    (Relu2): ReLU()
    (MaxPool2): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mo
  )
  (fully_connected): Linear(in_features=8192, out_features=16, bias=True)
)

```

▼ Training the Method1 Model

```

##Training the Method1 Model
start_time = time.time()
n_epochs = 25
for epoch in range(n_epochs):
    train_model(trainloader_small, epoch, model_ver2, False, False)

print('Finished Training')
print('\033[1m' + 'Time Taken to Train the Model: ' + '--- %s mins ---' % round((time.

```

```

☞ Finished Training
Time Taken to Train the Model: --- 0.54 mins ---

```

▼ Evaluating the Method1 Model on Test Data

```

test_model(testloader_small, model_ver2)

☞ Accuracy of the network on the test images: 64 %
Time Taken for Evaluating Model on Test: --- 0.06 seconds ---

```

We observe that Adding Batch Normalization Increases our Accuracy Marginally

▼ Method2: Adding Dropout Layer:

```

class CNN_Ver3(nn.Module):
    def __init__(self, input_channels=1, num_classes=16):
        super(CNN_Ver3, self).__init__()
        self.layer1 = nn.Sequential(

```

```

self.layer1 = nn.Sequential()
self.layer1.add_module("Conv1", nn.Conv2d(in_channels=input_channels, out_channels=16, kernel_size=3, stride=1, padding=1))
self.layer1.add_module("BN1", nn.BatchNorm2d(num_features=16))
self.layer1.add_module("Relu1", nn.ReLU())
self.layer1.add_module("MaxPool1", nn.MaxPool2d(kernel_size=2))

self.layer2 = nn.Sequential()
self.layer2.add_module("Conv2", nn.Conv2d(in_channels=16, out_channels=32, kernel_size=3, stride=1, padding=1))
self.layer2.add_module("BN2", nn.BatchNorm2d(num_features=32))
self.layer2.add_module("Relu2", nn.ReLU())
self.layer2.add_module("MaxPool2", nn.MaxPool2d(kernel_size=2))
self.layer2.add_module("dropout", nn.Dropout2d(0.001))

self.fully_connected = nn.Linear(32 * 16 * 16, num_classes)

def forward(self, x):
    x = self.layer1(x)
    x = self.layer2(x)
    x = x.view(x.size(0), -1)
    x = self.fully_connected(x)
    return x

# defining the model
model_ver3 = CNN_Ver3()
optimizer = optim.Adam(model_ver3.parameters(), lr=0.005)
criterion = nn.CrossEntropyLoss()

if torch.cuda.is_available():
    model_ver3 = model_ver3.cuda()
    criterion = criterion.cuda()

##Training the Method1 Model
start_time = time.time()
n_epochs = 25
for epoch in range(n_epochs):
    train_model(trainloader_small, epoch, model_ver3, False, False)

print('Finished Training')
print('\033[1m' + 'Time Taken to Train the Model: ' + '--- %s mins ---' % round((time.time() - start_time) / 60, 2))

test_model(testloader_small, model_ver3)

☞ Finished Training
Time Taken to Train the Model: --- 0.55 mins ---
Accuracy of the network on the test images: 62 %
Time Taken for Evaluating Model on Test: --- 0.04 seconds ---

```

We observe that Adding Dropout Layer doesn't Increase our Accuracy much, by

▼ Method3: Changing The Activation Function to Sigmoid and Leaky Relu:

```

class CNN_Ver4(nn.Module):
    def __init__(self, input_channels=1, num_classes=16):
        super(CNN_Ver4, self).__init__()
        self.layer1 = nn.Sequential()
        self.layer1.add_module("Conv1", nn.Conv2d(in_channels=input_channels, out_channels=16, kernel_size=3, padding=1))
        self.layer1.add_module("BN1", nn.BatchNorm2d(num_features=16))
        self.layer1.add_module("Sigmoid1", nn.Sigmoid())
        self.layer1.add_module("MaxPool1", nn.MaxPool2d(kernel_size=2))

        self.layer2 = nn.Sequential()
        self.layer2.add_module("Conv2", nn.Conv2d(in_channels=16, out_channels=32, kernel_size=3, padding=1))
        self.layer2.add_module("BN2", nn.BatchNorm2d(num_features=32))
        self.layer2.add_module("Sigmoid2", nn.Sigmoid())
        self.layer2.add_module("MaxPool2", nn.MaxPool2d(kernel_size=2))

        self.fully_connected = nn.Linear(32 * 16 * 16, num_classes)

    def forward(self, x):
        x = self.layer1(x)
        x = self.layer2(x)
        x = x.view(x.size(0), -1)
        x = self.fully_connected(x)
        return x

# defining the model
model_ver4 = CNN_Ver4()
optimizer = optim.Adam(model_ver4.parameters(), lr=0.001)
criterion = nn.CrossEntropyLoss()

if torch.cuda.is_available():
    model_ver4 = model_ver4.cuda()
    criterion = criterion.cuda()

##Training the Method1 Model
start_time = time.time()
n_epochs = 25
for epoch in range(n_epochs):
    train_model(trainloader_small, epoch, model_ver4, False, False)

print('Finished Training')
print('\033[1m' + 'Time Taken to Train the Model: ' + '--- %s mins ---' % round((time.time() - start_time) / 60))

test_model(testloader_small, model_ver4)

```



We observe that Changing Activation Function to *Sigmoid* deteriorates our Accuracy

▼ Let's check with Leaky Relu

```
class CNN_Ver5(nn.Module):
    def __init__(self, input_channels=1, num_classes=16):
        super(CNN_Ver5, self).__init__()
        self.layer1 = nn.Sequential()
        self.layer1.add_module("Conv1", nn.Conv2d(input_channels=input_channels, out_channels=16, kernel_size=3, padding=1))
        self.layer1.add_module("BN1", nn.BatchNorm2d(num_features=16))
        self.layer1.add_module("Lrelu1", nn.LeakyReLU())
        self.layer1.add_module("MaxPool1", nn.MaxPool2d(kernel_size=2))

        self.layer2 = nn.Sequential()
        self.layer2.add_module("Conv2", nn.Conv2d(input_channels=16, out_channels=32, kernel_size=3, padding=1))
        self.layer2.add_module("BN2", nn.BatchNorm2d(num_features=32))
        self.layer2.add_module("Lrelu2", nn.LeakyReLU())
        self.layer2.add_module("MaxPool2", nn.MaxPool2d(kernel_size=2))

        self.fully_connected = nn.Linear(32 * 16 * 16, num_classes)

    def forward(self, x):
        x = self.layer1(x)
        x = self.layer2(x)
        x = x.view(x.size(0), -1)
        x = self.fully_connected(x)
        return x

# defining the model
model_ver5 = CNN_Ver5()
optimizer = optim.Adam(model_ver5.parameters(), lr=0.001)
criterion = nn.CrossEntropyLoss()

if torch.cuda.is_available():
    model_ver5 = model_ver5.cuda()
    criterion = criterion.cuda()

##Training the Method1 Model
start_time = time.time()
n_epochs = 25
for epoch in range(n_epochs):
    train_model(trainloader_small, epoch, model_ver5, False, False)

print('Finished Training')
```

```
print('\033[1m' + 'Time Taken to Train the Model: ' + '--- %s mins ---' % round((time.

test_model(testloader_small,model_ver5)
```



We observe that Changing Activation Function to *LeakyRelu* gives more or less sa

▼ Method4: Augmenting Data by rotating the Image and checking Accuracy on Tra

```
# load data into size (64, 64)

start_time = time.time()
img_size = (64, 64)
batch_num = 50 # training sample number per batch

# load training dataset
trainloader_small_rot = list(load_dataset('./data/train/', img_size, batch_num=batch_n
                                augment=True, zero_centered=True, rotate_90=True)
train_num = len(trainloader_small_rot)
print("Finish loading %d minibatches(=%d) of training samples." % (train_num, batch_n

print('\033[1m' + 'Time Taken to load the Training: ' + '--- %s minutes ---' % round(
```



```
# # # Save intermediate image data into disk
# file = open('trainloader_small_rot.pkl','wb')
# pickle.dump(trainloader_small_rot, file)

# Load intermediate image data from disk
file = open('trainloader_small_rot.pkl', 'rb')
trainloader_small_rot = pickle.load(file)
file.close()
# Verify number of training samples after rotating images and augmenting
print('\033[1m' + 'Training Samples: ' + str(len(trainloader_small_rot)*50))
```



```
## Training using the CNN_Ver1, model_ver1

start_time = time.time()
# defining the number of epochs
n_epochs = 25
# empty list to store training losses
train_losses = []
# loop over the dataset multiple times
for epoch in range(n_epochs):
    train_model(trainloader_small_rot,epoch,model_ver1, False, True)

print('Finished Training')
print('\033[1m' + 'Time Taken to Train the Model: ' + '--- %s mins ---' % round((time.
```



```
# plotting the training loss
plt.plot(train_losses, label='Training loss')
plt.title('Loss on the Augmented Training Set using Image Rotation',size= 13)
plt.legend()
plt.xlabel('Epochs',size= 15)
plt.ylabel('Loss',size=15)
```

```
plt.show()
```



```
file = open('testloader_small.pkl', 'rb')
testloader_small = pickle.load(file)
file.close()
test_model(testloader_small,model_ver1)
```



We observe that with Rotating Image and Augmentating in the Training Set, our accuracy falls by a significance

Complete Report for the Question 1:

Part1 :

- On the given 2400 training and 400 test images, we augmented Data by doing mirror image
- We Normalized the training and test data, and centered around zero
- Build first CNN model with below parameters:

Two Hidden Layers, [Kernel Size = 5 with Stride= 1 and padding =2], Activation Function : ReLU, stride=2, padding=0]

- Time taken for training the data set 0.48mins and time taken for evaluating the model on test data
- **Accuracy Observed = 62%**

Part2 : We Apply below Methods to further investigate and improve our Accuracy:

Method1: Batch Normalization


```
Layer1: nn.BatchNorm2d(num_features=16) , Layer2: nn.BatchNorm2d(num_features=32)
```

Accuracy Observed = 64%

Method2: Adding Dropout Layer

```
Layer2: nn.Dropout2d(0.001)
```

Accuracy Observed = 62%

Method3: Changing Activation Function to Sigmoid , Leaky Relu

Adding Sigmoid and Leaky Relu in Layer1 and Layer2:

Accuracy Observed = 58% with nn.Sigmoid()

Accuracy Observed = 64% with nn.LeakyReLU()

Method4: Image Rotation

Rotating Images by 90degree and augmenting in the dataset, our 4800 Training Set doubles to 9600

Accuracy Observed = 36%, worst performance observed

Summary : Best Accuracy observed = 64% by using RELU Activation Function and Batch Normalisation

▼ Problem 2: Fine Tuning a Pre-Trained Deep Network

{Part 1: 30 points} Our convolutional network to this point isn't "deep". Fortunately, the representations learned by the network are such that they generalize surprisingly well to other recognition tasks.

But how do we use an existing deep network for a new recognition task? Take for instance, [AlexNet](#) network corresponding to 1000 ImageNet categories.

Strategy A: One could use those 1000 activations as a feature in place of a hand crafted feature such as SIFT. One would train a classifier (typically a linear SVM) in that 1000 dimensional feature space. However, those features are very specific and may not generalize well to new recognition tasks. It is generally better to use the activations from a deeper layer, e.g. the 4096 activations in the last 2nd fully-connected layer. You can often get away with sub-sampling the activations, e.g. taking only the first 200 activations.

Strategy B: *Fine-tune* an existing network. In this scenario you take an existing network, replace the first few layers and train the entire network again with images and ground truth labels for your recognition task. You use the pre-trained network as a better initialization than the random weights used when training from scratch. When you have a simple network from scratch (e.g. with the 16 classes) this is an attractive option. Fine-tuning can warm-start the network's activations directly from an pre-trained CNN. For example, in [this paper](#) from CVPR 2015, there wasn't training from scratch, but fine tuning led to 4 times higher accuracy than using off-the-shelf networks directly.

You are required to implement **Strategy B** to fine-tune a pre-trained **AlexNet** for this scene classification task to achieve a performance of 85% approximately. It takes roughly 35~40 minutes to train 20 epoches with AlexNet.

Please provide detailed descriptions of:

- (1) which layers of AlexNet have been replaced
- (2) the architecture of the new layers added including activation methods (same as problem 1)
- (3) the final accuracy on test set along with time consumption for both training and testing

{Part 2: 20 points} Implement Strategy A where you use the activations of the pre-trained network as features for the scene classification task. Report the final accuracy on test set along with time consumption for both training and testing.

{Bonus: 10 points} Bonus will be given to those who fine-tune the [VGG network paper](#) and compare performance with the AlexNet. Did it performed better or worse.

Hints:

- Many pre-trained models are available in PyTorch at [here](#).
- For fine-tuning pretrained network using PyTorch, please read this [tutorial](#).

```
# reload data with a larger size
img_size = (224, 224)
batch_num = 50 # training sample number per batch

# load training dataset
trainloader_large = list(load_dataset('./data/train/', img_size, batch_num=batch_num,
                                     augment=False, is_color=True, zero_centered=True))
train_num = len(trainloader_large)
print("Finish loading %d minibatches(=%d) of training samples." % (train_num, batch_num))

# load testing dataset
testloader_large = list(load_dataset('./data/test/', img_size, num_per_class=50, batch_num=batch_num))
test_num = len(testloader_large)
print("Finish loading %d minibatches(=%d) of testing samples." % (test_num, batch_num))

# file = open('trainloader_large.pkl', 'wb')
# pickle.dump(trainloader_large, file)
# file.close()

# file = open('testloader_large.pkl', 'wb')
```

```
# pickle.dump(testloader_large, file)
# file.close()

# Load intermediate image data from disk
file = open('trainloader_large.pkl', 'rb')
trainloader_large = pickle.load(file)
file.close()

file = open('testloader_large.pkl', 'rb')
testloader_large = pickle.load(file)
file.close()

print('\033[1m' + 'Total Large Size Training Samples: ' + str(len(trainloader_large)*50))
print('\033[1m' + 'Total Large Size Test Samples: ' + str(len(testloader_large)*50))

print('\033[1m' + 'Total Large Size Training Batches: ' + str(len(trainloader_large)))
print('\033[1m' + 'Total Large Size Test Batches: ' + str(len(testloader_large)))
```



Strategy B : Fine Tuning Alex Net

```
# =====
#         Fine-Tune Pretrained Alex Network
# =====

num_classes=16
import torchvision.models as models

##Initialising the Model for the Alex Net for fine tuning with the ouput Layer to 16 classes
def initialize_model(num_classes, use_pretrained=True):
    model_ft = models.alexnet(pretrained=use_pretrained)
    num_fttrs = model_ft.classifier[6].in_features
    model_ft.classifier[6] = nn.Linear(num_fttrs,num_classes)
    return model_ft

# Initialize the model
model_ft = initialize_model(num_classes, use_pretrained=True)
print(model_ft)
```



▼ Create the Optimizer

```
#checking if GPU is available
if torch.cuda.is_available():
    model_ft = model_ft.cuda()

params_to_update = model_ft.parameters()

optimizer_ft = optim.SGD(params_to_update, lr=0.001, momentum=0.9)
criterion = nn.CrossEntropyLoss().cuda()

##Training the Method1 Model
start_time = time.time()
n_epochs = 25
train_losses=[]

for epoch in range(n_epochs):
    print('\033[1m' + 'Performance Statistics for Epoch Number :' + str(epoch + 1))
    train_running_loss = 0.0
    train_total_loss= 0.0
    for batch_num, data in enumerate(trainloader_large, 0):
        # getting the inputs; data is a tensor tuple list of [inputs, labels]
        x_train, y_train = data
        x_train = x_train.float()

        # converting the data into GPU format
        if torch.cuda.is_available():
            x_train = x_train.cuda()
            y_train = y_train.cuda()
```

```

optimizer_ft.zero_grad()
output_train = model_ft(x_train)
loss_train = criterion(output_train, y_train)
loss_train.backward()
optimizer_ft.step()

# print statistics
train_running_loss += loss_train.item()
train_total_loss += loss_train.item()

if batch_num % 20 == 19:    # print every 20 mini-batches
    # print('\033[1m' + '[Batch NO.: %d, Mini Batch: %5d] loss: %.3f' %
    #       (epoch + 1, batch_num + 1, train_running_loss / 20))
    train_running_loss = 0.0
train_losses.append(train_total_loss/(batch_num + 1))

print('\033[1m' + '[Avg Loss for the Epoch NO.: %d, Loss: %.3f]' % (epoch + 1, tra
print('\n')

print('Finished Training')
print('\033[1m' + 'Time Taken to Train the Model: ' + '--- %s mins ---' % round((time.

```



```
# plotting the training loss
plt.plot(train_losses, label='Training loss')
plt.title('Loss on the Alex Training Set on Successive Epochs',size= 13)
plt.legend()
plt.xlabel('Epochs',size= 15)
plt.ylabel('Loss',size=15)
plt.show()
```



```

# Evaluating Network
# =====

def test_model(test_data,model_version):
    start_time = time.time()
    # prediction and Accuracy for Test set
    correct = 0
    total = 0
    model_version.eval()
    for batch_num, data in enumerate(test_data, 0):
        x_test, y_test = data
        x_test = x_test.float()

        if torch.cuda.is_available():
            x_test = x_test.cuda()
            y_test = y_test.cuda()

        y_pred = model_version(x_test.cuda())
        _, predicted = torch.max(y_pred.data, 1)
        total += y_test.size(0)
        correct += (predicted == y_test).sum().item()

    print('\033[1m' + 'Accuracy of the network by tuning the Alex Network on the test
    print('\033[1m' + 'Time Taken for Evaluating Model on Test: ' + '--- %s seconds --

test_model(testloader_large,model_ft)

```



Strategy A: Extracting Features from the Activation functions of the 2nd last Layer of the Alex Net ar

```

model_feat = models.alexnet(pretrained=True)
## Removing the last layer and using the 2nd last layer activation functions output as
for param in model_feat.parameters():
    param.requires_grad = False

new_classifier = nn.Sequential(*list(model_feat.classifier.children())[:-1])
model_feat.classifier = new_classifier
print(model_feat)

```



```
## Defining Features Extraction Method
```

```
def alex_feat(model,data):
    model = model.cuda()
    start_time = time.time()
    feat_size= model.classifier[4].in_features
    alex_features = np.empty((0, feat_size))
    feat_labels=np.zeros((0,1))

    for batch_num, batch_data in enumerate(data, 0):
        x_train, y_train = batch_data
        x_train = x_train.float()
        y_train= np.reshape(y_train,(y_train.shape[0],1))
        y_train= y_train.numpy()

        x_train = x_train.cuda()
        features = model(x_train).cpu().data.numpy()
        features=features[:, :feat_size]
        alex_features= np.append(alex_features, features, axis=0)
        feat_labels=np.vstack((feat_labels,y_train))

    print('\033[1m' + 'Time Taken for Extracting Features: ' + '--- %s seconds ---'
    return np.asarray(alex_features.astype('float32')), np.asarray(feat_labels.astype('float32'))

x_train,y_train= alex_feat(model_feat,trainloader_large)
print('\033[1m' + 'Shape of the Training Feature Vector: ' + str(x_train.shape))
print('\033[1m' + 'Shape of the Training Label Vector: ' + str(y_train.shape))
print('\n')
x_test,y_test= alex_feat(model_feat,testloader_large)
print('\033[1m' + 'Shape of the Test Feature Vector: ' + str(x_test.shape))
print('\033[1m' + 'Shape of the Test Label Vector: ' + str(y_test.shape))
```




Training the SVM Classifier and Model Accuracy

```
# training a linear SVM classifier
start_time = time.time()
from sklearn import svm
from sklearn.svm import LinearSVC
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score

svm_model_linear = svm.LinearSVC(random_state=0,tol=1e-5,C=1,max_iter= 5000)
svm_model_linear.fit(x_train, y_train.ravel())

test_predictions = svm_model_linear.predict(x_test)
accuracy= accuracy_score(y_test, test_predictions)

print ('\033[1m' + 'Accuracy Observed with SVM Classifier = ' + str(accuracy*100 ) +
print ('\033[1m' + 'Time Taken for Classification: ' + '--- %s seconds ---' % round((ti
```



Accuracy observed by using 2nd Last Layer Activation Functions as our feature gives a highly improved accuracy of 78.5% compared to the previous Bag of Sift

Bonus: Fine Tuning VGG Network

```
# =====
# Fine-Tune Pretrained Alex Network
# =====
num_classes=16

import torchvision.models as models

##Initialising the Model for the Alex Net for fine tuning with the ouput Layer to 16 c

def initialize_model(num_classes, use_pretrained=True):
    model_ft = None
    input_size = 0
    model_ft = models.vgg16(pretrained=use_pretrained)
    num_ft_classes = model_ft.classifier[6].in_features
```

```
num_fttrs = model_ft.classifier[6].in_features
model_ft.classifier[6] = nn.Linear(num_fttrs,num_classes)

return model_ft

# Initialize the model
model_vgg = initialize_model(num_classes, use_pretrained=True)
print(model_vgg)
```



```
#checking if GPU is available
if torch.cuda.is_available():
    model_vgg = model_vgg.cuda()

feature_extract= False
params_to_update = model_vgg.parameters()
print("Params to learn:")

for name,param in model_vgg.named_parameters():
    if param.requires_grad == True:
        print("\t",name)

# All parameters are being optimized
optimizer_ft = optim.SGD(params_to_update, lr=0.001, momentum= 0.9)
exp_lr_scheduler= optim.lr_scheduler.StepLR(optimizer_ft,step_size=7,gamma= 0.1)

criterion = nn.CrossEntropyLoss()
```



```
###Training the Method1 Model
```

```
start_time = time.time()
```

```

start_time = time.time(),
n_epochs = 20
train_losses=[]

for epoch in range(n_epochs):
    model_vgg.train()
    print('\033[1m' + 'Performance Statistics for Epoch Number :' + str(epoch + 1))
    train_running_loss = 0.0
    train_total_loss= 0.0
    for batch_num, data in enumerate(trainloader_large, 0):
        # getting the inputs; data is a tensor tuple list of [inputs, labels]
        x_train, y_train = data
        x_train = x_train.float()

        # converting the data into GPU format
        if torch.cuda.is_available():
            x_train = x_train.cuda()
            y_train = y_train.cuda()

        optimizer_ft.zero_grad()
        output_train = model_vgg(x_train)
        loss_train = criterion(output_train, y_train)
        loss_train.backward()
        optimizer_ft.step()

        # print statistics
        train_running_loss += loss_train.item()
        train_total_loss += loss_train.item()

        if batch_num % 20 == 19:    # print every 20 mini-batches
            print('\033[1m' + '[Batch NO.: %d, Mini Batch: %5d] loss: %.3f' %
                  (epoch + 1, batch_num + 1, train_running_loss / 20))
            train_running_loss = 0.0
    train_losses.append(train_total_loss/(batch_num + 1))

    print('\033[1m' + '[Avg Loss for the Epoch NO.: %d, Loss: %.3f]' % (epoch + 1, train_total_loss / (epoch + 1)))
    print('\n')

print('Finished Training')
print('\033[1m' + 'Time Taken to Train the Model: ' + '--- %s mins ---' % round((time.time() - start_time) / 60))

```



```
# =====  
#           Evaluating Network  
# =====
```

```
def test_model(test_data,model_version):  
    start_time = time.time()  
    # prediction and Accuracy for Test set  
    correct = 0  
    total = 0
```

```

model_version.eval()
for batch_num, data in enumerate(test_data, 0):
    x_test, y_test = data
    x_test = x_test.float()

    if torch.cuda.is_available():
        x_test = x_test.cuda()
        y_test = y_test.cuda()

    y_pred = model_version(x_test.cuda())
    _, predicted = torch.max(y_pred.data, 1)
    total += y_test.size(0)
    correct += (predicted == y_test).sum().item()

print('\033[1m' + 'Accuracy of the network by tuning the VGG Network on the test i
print('\033[1m' + 'Time Taken for Evaluating Model on Test: ' + '--- %s seconds --

test_model(testloader_large,model_vgg)

```



Tuning the Network with Alex Net increases our accuracy by 4-5%

Complete Report for the Question 2:

Part1 : Implementing Startegy B --> Fine tuning Alex Net

Changing the Number of output features in the final layer of the classifier to 16: Linear(in_feature bias=True)

- Time taken for training the New Alex Net with our Training data set= 2.2mins and time taken test data 0.36secs
- **Accuracy Observed = 80%**

Part2 : Implementing Startegy A --> Using Activation Functions of 2nd Last Layer as feature vector a

Extracting Features from the Activation functions of the 2nd last Layer of the Alex Net and using :

- Linear(in_features=4096, out_features=4096, bias=True), ReLU(inplace=True)
- Time taken for extracting the training features form the Alex penultimate layer= 2.55secs and test features = 0.34secs
- Time taken for SVM classification = 2.88secs
- **Accuracy Observed = 78.5%, highly improved accuracy comapred with Bag of SIFT descrip assignemnt**

Part3 : Bonus --> Fine tuning VGG Net

Colab has a good feature of version control, you should take advantage of this to save your work properly. The submission made in blackboard is the only one that we consider for grading. To be more specific, we consider the submission made right before the timestamp of the submission made in blackboard.

You are encouraged to post and answer questions on Piazza. Based on the amount of email that we handle, we dealys in replying to personal emails. Please ask questions on Piazza and send emails only for personal matters. Be aware that your code will undergo plagiarism check both vertically and horizontally. Please do your

Late submission penalty:

There will be a 10% penalty per day for late submission. However, you will have 4 days throughout the semester for late submission. Note that the grace period is calculated by days instead of hours. If you submit the homework after the deadline, the 10% penalty will be applied.